

Report-ASSIGNMENT 2 OS – (CE21BTECH11008)

System details:- Intel core I5 , 8 logical cores.

Reading The file My input file name is "inp.txt" and output file name is "out.txt" . For OutFile I have made a file pointer that is called at the end of the program and getting written using answer matrix. For input file using ifstream from fstream library and taking input in the main () function. "In.txt" contains a line. Extracting the N and K from this by simple reading and called "read matrix" function From reading determining the N and K value and making "matrix" 2d vector of N x N size and storing the value of K and N using pass by reference.

To measure Time:

Used "auto start_time = std::chrono::high_resolution_clock::now();" found in time.h library . As we are intended to find the time taken by threads to execute on different set of assigned rows. So, I am starting the time for each thread when it is entering in that function and end_time when that thread is just about to call return. Taking difference to find interval. Now made one vector name "time" that is keeping track record of the time taken by each thread. It is free from race conditions because every thread has its unique index. Finally, taking the average of the time taken by each thread.

Threading Methodology

I have used Posix thread and called pthread_create function with its arguments and for parameter passing I have made struct data type. That is containing N and K and 2d vector matrix to be passed as this is the data used by all. For passing pthread_create I have made array of pthread_t of size k where the function instance will be stored. Also, I am passing ans matrix that is keeping track of the result. And last have called pthread_join function as no return value is there so passed its argument as null and thread id th[i].

for load balancing

For Chunk: - same strategy applying as guided in the question like finding ceiling(N/K) and calling as Dist and Distribution each thread this amount of row in consecutive manner. Ofcourse there is chances that Last thread will lesser number of rows, but it will not affect much.

For Mixed: - As the question suggests, I have made a 2d vector for this that will keep the record of row assigned to each thread. This is done using the main thread. I have done its implementation using modulo function. Like just travelling in each row index and assign to thread p if $p = \text{row} \% K$ (i.e. row modulo thread size).

Why chunk is not always a good approach.

see a triangular matrix. if the things will be provided in chunk to the thread, then last threads literally have very few works to do (if it is upper triangular) as thread will get 0 value and sit idle while rest of thread will have comparatively higher value. It may induce load imbalance if chunks have significantly different computational loads. But it favors locality principal of the cache thus may improve some efficiency if matrix is not biased like triangular or other thus mixed way of matrix multiplication provide better efficiency in this case. Matrices with specific patterns often arise in structured problems. For example, in signal processing or communications, matrices with zeros coming at regular interval is common in certain rows or columns representing inactive or irrelevant components thus in this kind of problem mixed way will be inefficient.

Why is mixed way of matrix multiplication also not very efficient? Even though it encounters the drawback of the chunk but it does not able to recover the best property of chunk that chunk provide locality. see for larger number of threads say like 8 to 10. Each row needed to do this multiplication will be k rows apart. But see in the cache perspective cash can load few rows due to locality of cache behavior (say taking 4 rows). But as in this methodology we using 8th row after second iteration thus most of the time lead to cache miss and thus may lead to higher time

Setting affinity: -

Approach 1(Not full proof): -

First creating the thread and then using “pthread_setaffinity_np()” function to assign the fixed core.

But problem with this is when thread has already been created and function is passed then it will start executing on different thread. At the beginning this will assign to some different thread but when “pthread_setaffinity_np()” is called then it will assign to particular desired core.

```
pthread_create(&thread, NULL, threadFunc, (void *)threadId);  
// Set the CPU affinity for this thread  
cpu_set_t cpuset;  
CPU_ZERO(&cpuset);  
CPU_COUNT(&cpuset);  
CPU_SET(i % 8, &cpuset);  
  
pthread_setaffinity_np(thread, sizeof(cpu_set_t), &cpuset);
```

my_experiment: - I have tried to simulate . When it is entering in the function I have printed on which cpu it is running using “shed_getcpu()” function. In result I got that most of the time it is getting assigned to particular core but some time it is missing due to the rason I have mentioned above.

Thus this is not a full proof method.

```
void *threadFunc(void *arg)
{
    //restoring the argument that are passed
    int threadId = *(int *)arg;

    // Get the CPU core on which this thread is running
    int cpu = sched_getcpu();

    std::cout << "Thread " << threadId << " is running on CPU " << cpu
    << std::endl;

    // Do some work here
    usleep(1000); // Sleep for 1 second

    return NULL;
}
```

Approach(full proof)

This is the approach I have used. First setting the affinity part of attribute to fixed core using the function “cpu_attr_setaffinity_np()” and then creating the thread using “pthread_create” with the same attribute passing through this create function . This will ensure that thread from beginning itself is assigned to the fixed thread.

```
pthread_t thread;
cpu_set_t cpuset;
CPU_ZERO(&cpuset);
int *threadId = new int(i);

CPU_SET(*threadId % nc, &cpuset);
pthread_attr_t attr;
//assigning the attribute a fixed core form cpuset
pthread_attr_setaffinity_np(&attr, sizeof(cpu_set_t),
&cpuset);

//passing that attribute when creating the thraed
pthread_create(&thread, &attr, threadFunc, (void *)threadId);
```

Using above simulation in previous approaches I have checked multiple times that it is getting assigned to fixed core.

Determination of Cache miss

Main motivation of using affinity is to reduce the cache miss because if a thread will be bounded to a particular cpu then it can use previous loaded data in cache and cost of invalidation of data will less.

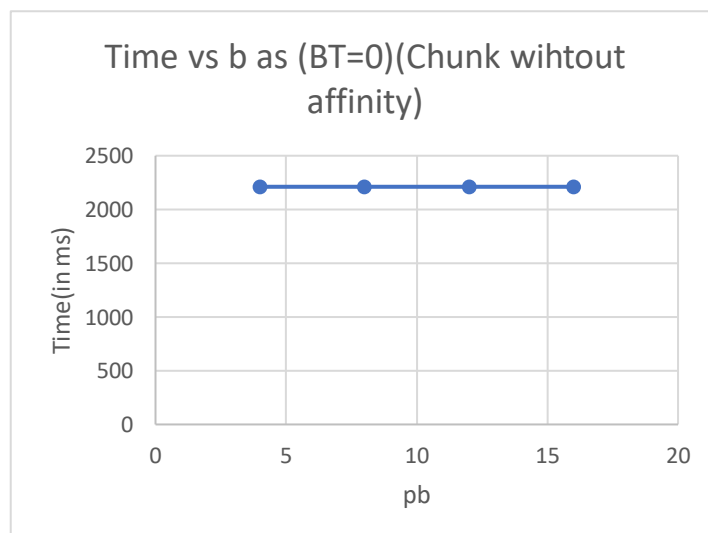
For this I have added additional linux performance tool and used perf command to

“perf stat -cache-misses, cache-references ./executable_file” this will provide the total %cache missed in this code.

Experiment 1


Part (1) Chunk algorithm without threads being bound to cores

| Chunk with 8 cores and 32 threads with No of Bounded threads | | | | | | | | | | | | | | | | |
|--|-----|----|-------------|-------|-------|-------|-------|------------------|--|------------|-------|-------|-------|-------|-----------|--|
| | | | time(in ms) | | | | | average time(ms) | | cache miss | | | | | Avg(in %) | |
| | b=4 | BT | obs-1 | obs-2 | obs-3 | obs-4 | obs-5 | | | obs-1 | obs-2 | obs-3 | obs-4 | obs-5 | | |
| b | 4 | 0 | 2226 | 2258 | 2178 | 2254 | 2141 | 2211.4 | | 4.97 | 4.09 | 5.28 | 3.86 | 4.15 | 4.47 | |
| 2b | 8 | 0 | 2226 | 2258 | 2178 | 2254 | 2141 | 2211.4 | | 4.97 | 4.09 | 5.28 | 3.86 | 4.15 | 4.47 | |
| 3b | 12 | 0 | 2226 | 2258 | 2178 | 2254 | 2141 | 2211.4 | | 4.97 | 4.09 | 5.28 | 3.86 | 4.15 | 4.47 | |
| 4b | 16 | 0 | 2226 | 2258 | 2178 | 2254 | 2141 | 2211.4 | | 4.97 | 4.09 | 5.28 | 3.86 | 4.15 | 4.47 | |



Part(B) Chunk algorithm with a given b

| Chunk with 8 cores and 32 threads with BT number of Bounded threads | | | | | | | | | | | | | | | |
|---|-----|-------------|-------|-------|-------|-------|------------------|--------|------------|-------|-------|-------|-------|-----------|--|
| | | time(in ms) | | | | | average time(ms) | | cache miss | | | | | Avg(in %) | |
| | b=4 | BT(b=4) | obs-1 | obs-2 | obs-3 | obs-4 | obs-5 | | obs-1 | obs-2 | obs-3 | obs-4 | obs-5 | | |
| b | 4 | 8 | 2075 | 2092 | 2100 | 2075 | 2047 | 2077.8 | 3.36 | 3.44 | 2.78 | 3.01 | 3.63 | 3.244 | |
| 2b | 8 | 16 | 2112 | 2101 | 2080 | 2100 | 2029 | 2084.4 | 3.44 | 3.27 | 3.39 | 3.56 | 3.64 | 3.46 | |
| 3b | 12 | 24 | 2151 | 2111 | 2149 | 2062 | 2085 | 2111.6 | 3.52 | 3.61 | 3.55 | 3.54 | 3.42 | 3.528 | |
| 4b | 16 | 32 | 2118 | 2091 | 2071 | 2022 | 2038 | 2068 | 3.42 | 3.45 | 3.25 | 3.86 | 3.5 | 3.496 | |



Individual observation

In part A

As no bounding is there time will be same.

Note: for the graph we have to draw between time vs Number of BT. But in this case BT = 0. I have changed the X axis with pb value where $b = K/C$

In part B.

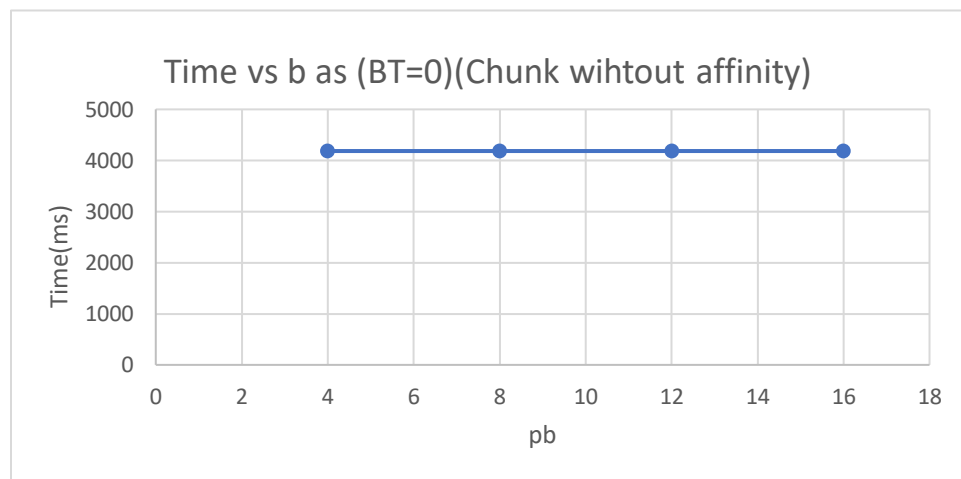
It is observed that increasing the BT will not necessarily increase performance. As can be seen when bounded thread increases 8 to 24 then time is increasing. This may be due to the increasing the bounded thread more than core may lead to condition like even one core has completed their work while other has have piling up of threads. In this case load balancing by scheduler is not possible thus leading to no utilization of some core and thus time increases.

Overall Observations

Chunk with affinity set is doing better than chunk with affinity set. Average time of part B is approximately 100ms less than the average time of part A. Also, a significant reduction in cache miss can be seen when CPU affinity is applied. Thus, bounding the thread in this case is significantly improving the performance in the execution time and cache miss. Also, it is observed that execution time is not always proportional to cache miss. This may be due to the fact that execution time depends on many other factors like vacancy of processor, current speed of processor, cache miss etc. Thus, cache miss is one of the factors in deciding this.

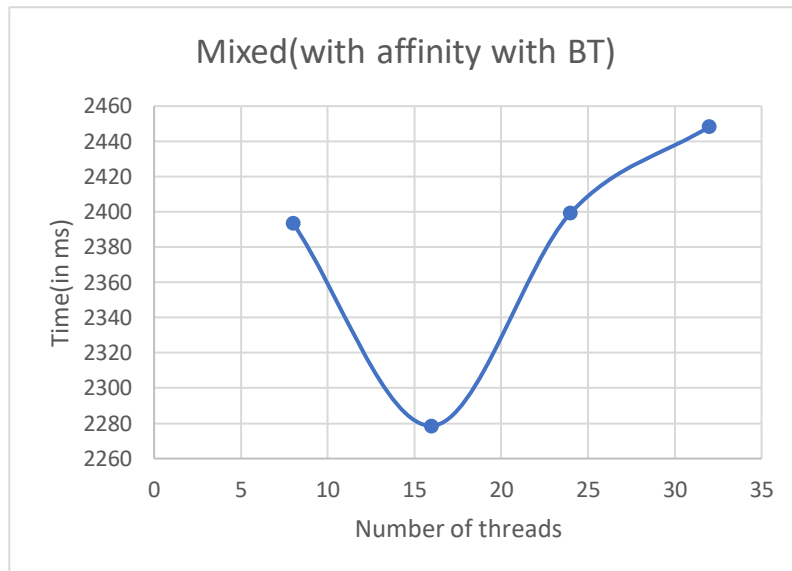
Part (3) (Mixed algorithm without threads being bound to cores)

| Mixed with 8 cores and 32 threads with No of Bounded threads | | | | | | | | | | | | | | | |
|--|-----|----|-------------|-------|-------|-------|-------|------------------|------------|-------|-------|-------|-------|-------|--|
| | b=4 | BT | time(in ms) | | | | | average time(ms) | cache miss | | | | | in % | |
| | | | obs-1 | obs-2 | obs-3 | obs-4 | obs-5 | | obs-1 | obs-2 | obs-3 | obs-4 | obs-5 | | |
| b | 4 | 0 | 4233 | 4180 | 4179 | 4137 | 4195 | 4184.8 | 3.14 | 3.4 | 3.23 | 3.87 | 3.07 | 3.342 | |
| 2b | 8 | 0 | 4233 | 4180 | 4179 | 4137 | 4195 | 4184.8 | 3.14 | 3.4 | 3.23 | 3.87 | 3.07 | 3.342 | |
| 3b | 12 | 0 | 4233 | 4180 | 4179 | 4137 | 4195 | 4184.8 | 3.14 | 3.4 | 3.23 | 3.87 | 3.07 | 3.342 | |
| 4b | 16 | 0 | 4233 | 4180 | 4179 | 4137 | 4195 | 4184.8 | 3.14 | 3.4 | 3.23 | 3.87 | 3.07 | 3.342 | |



Part(4) Mixed algorithm without threads being bound to cores

| Mixed with 8 cores and 32 threads with BT number of Bounded threads | | | | | | | | | | | | | | | |
|---|-----|---------|-------------|-------|-------|-------|-------|------------------|------------|-------|-------|-------|-------|-----------|--|
| | b=4 | BT(b=4) | time(in ms) | | | | | average time(ms) | cache miss | | | | | Avg(in %) | |
| | | | obs-1 | obs-2 | obs-3 | obs-4 | obs-5 | | obs-1 | obs-2 | obs-3 | obs-4 | obs-5 | | |
| b | 4 | 8 | 2404 | 2380 | 2350 | 2422 | 2412 | 2393.6 | 3.36 | 3.32 | 3.23 | 3.33 | 3.39 | 3.326 | |
| 2b | 8 | 16 | 2397 | 2330 | 2309 | 2328 | 2029 | 2278.6 | 3.99 | 3.4 | 3.85 | 3.56 | 3.66 | 3.692 | |
| 3b | 12 | 24 | 2432 | 2394 | 2448 | 2349 | 2374 | 2399.4 | 3.52 | 3.37 | 3.52 | 3.76 | 3.62 | 3.558 | |
| 4b | 16 | 32 | 2463 | 2459 | 2489 | 2396 | 2435 | 2448.4 | 4.11 | 4.06 | 3.29 | 3.45 | 3.83 | 3.748 | |



Individual Observation

Part C. This is the simple mixed methodology used without any bounding. Its time will be constant as $BT = 0$ (always constant). The time taken by this method is quite high even though cash miss rate is less. It means other dependencies are playing a role in the determination of this time.

Part-D

It provides better timing than mixed without affinity. As the bounded thread is increasing the time first decreases then increases. This may be an initial increase of bounded thread gives guarantee that this thread will at least be there and thus decreasing the chances to make the CPU idle as some sort of assurance is coming. Also, increasing time after 16 threads can be due to overhead created by this thread on some core. Cache might have less work but some process on other core will be starving as they are bonded with those cores. Thus, this may be one reason for that.

Overall comparison: - Using affinity reasonably increases performance and decreases the time by 1.5 seconds on my processor. Cache miss doesn't have a good record in comparison to when it is unbounded.

Performance


Approximate conclusion can be drawn from the above experiment.

Performance of (Chunk with bounded thread) > (chunk without bounded thread) > (Mixed with bounded thread) > (Mixed with bounded thread).

Experiment 2


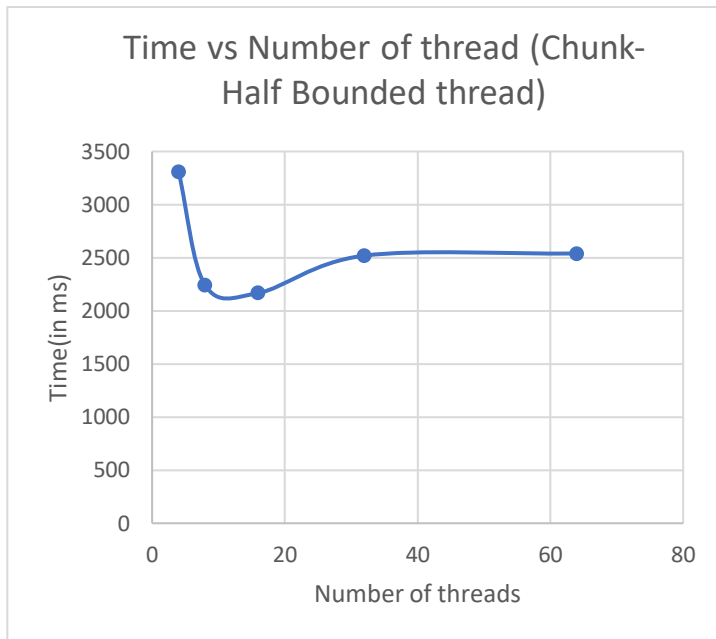
Part 5. – Average time execution without threads being bound to cores

Time vs Number of threads (without any CPU affinity)



(Part 6 & Part 7) Time vs Number of threads (without any CPU affinity) -first set is bounded 2nd is set unbounded

| | | Time vs Number of thread(Half Chunk - Bounded & Unbounded) | | | | | | | | | | | | | |
|------------------|------------------|--|------|------|------|------|------|------------------|--|------------|------|------|------|------|-------|
| | | time(in ms) | | | | | | average time(ms) | | cache miss | | in % | | | |
| Bounded | number of thread | | | | | | | | | | | | | | |
| (BT) | 2 | 4 | 3233 | 3307 | 3303 | 3310 | 3383 | 3307.2 | | 4.07 | 3.97 | 4.39 | 3.87 | 3.86 | 3.91 |
| (BT) | 4 | 8 | 2196 | 2241 | 2314 | 2297 | 2152 | 2240 | | 3.19 | 3.15 | 4.71 | 4.11 | 3.65 | 3.762 |
| (BT) | 8 | 16 | 2133 | 2181 | 2203 | 2211 | 2110 | 2167.6 | | 3.98 | 3.8 | 3.49 | 3 | 3.57 | 3.568 |
| (BT) | 16 | 32 | 2350 | 2568 | 2534 | 2577 | 2569 | 2519.6 | | 4.74 | 4.7 | 3.59 | 3.96 | 3.05 | 4.008 |
| (BT) | 32 | 64 | 2563 | 2533 | 2540 | 2549 | 2512 | 2539.4 | | 4.88 | 3.36 | 3.65 | 3.6 | 4.88 | 4.074 |
| | | | | | | | | | | | | | | | |
| Unbounded thread | | | | | | | | | | | | | | | |
| | 2 | 4 | 1719 | 1765 | 1777 | 1758 | 1815 | 1766.8 | | 4.07 | 3.97 | 4.39 | 3.87 | 3.86 | 3.91 |
| | 4 | 8 | 2197 | 2238 | 2311 | 2285 | 2150 | 2236.2 | | 3.19 | 3.15 | 4.71 | 4.11 | 3.65 | 3.762 |
| | 8 | 16 | 2085 | 2109 | 2039 | 2095 | 1938 | 2053.2 | | 3.98 | 3.8 | 3.49 | 3 | 3.57 | 3.568 |
| | 16 | 32 | 2273 | 2293 | 2160 | 2175 | 2158 | 2211.8 | | 4.74 | 4.7 | 3.59 | 3.96 | 3.05 | 4.008 |
| | 32 | 64 | 2098 | 2186 | 2017 | 2120 | 2125 | 2109.2 | | 4.88 | 3.36 | 3.65 | 3.6 | 4.88 | 4.074 |




Observation --- Here when comparing between bounded and unbounded the result came is astonishing. In this case the set of threads which is bounded is taking more than the threads which are unbounded. There is approximately 500 ms difference between their time. This may be due to internal inefficiency associated and due to overhead on scheduler. In this case if some bounded thread take more time then it will starve the thread waiting on same cpu as they can run only after this. This may lead to starvation of few threads in bounded one leading to high time than Normal threads as somehow scheduler optimises the thread at its end.

other observation – As the number of threads is increasing the is saturating. It may be due to we have created enough number of bounded and normal thread than the capacity of processor at which it can process.

Also when we compare the cache miss rate an important observation came out :

Approx – cache miss rate (all thread are bounded) cache miss rate(half bound and half normal)
 < cache miss rate (all thread are normal)


(Part 8) Without any bounding (as in first experiment)



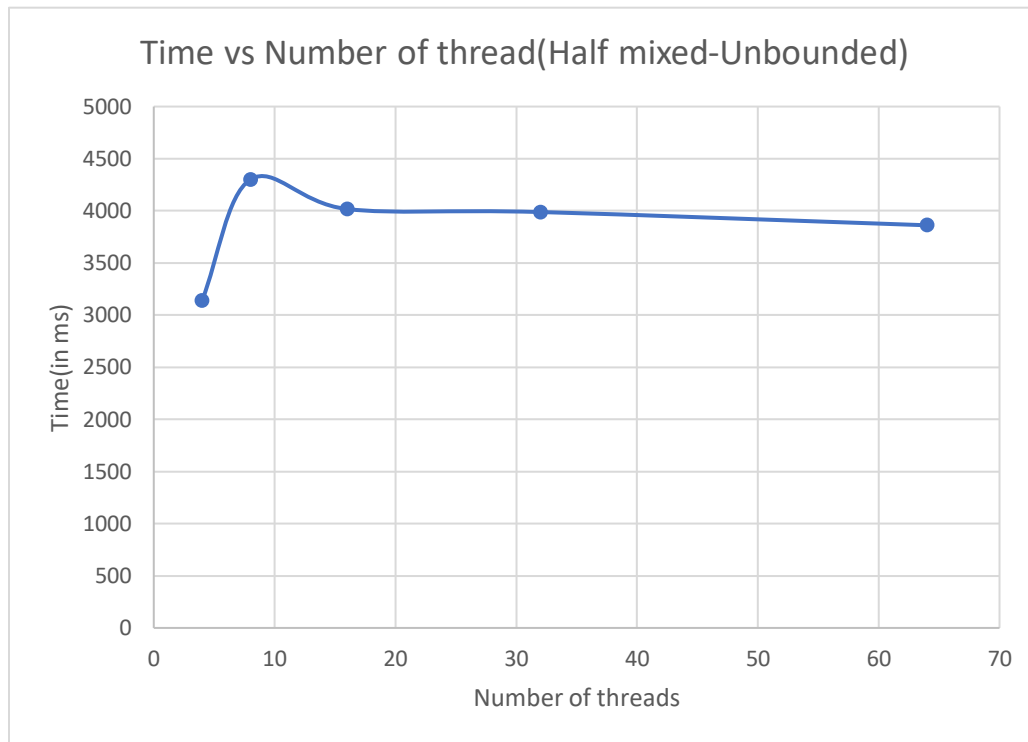
Part 9 & 10 (half thread bounded –first table and other half unbounded – second table)

| | | Time vs Number of thread(Half Mixed - Bounded & Unbounded) | | | | | | | | | | | | | |
|------|----|--|------|------|------|------------------|------|--------|--|------------|------|------|------|----------|-------|
| | | time(in ms) | | | | average time(ms) | | | | cache miss | | | | Avg in % | |
| (BT) | 2 | 4 | 5300 | 5288 | 5314 | 5224 | 5262 | 5277.6 | | 4.8 | 3.88 | 5 | 3.78 | 4.12 | 4.316 |
| (BT) | 4 | 8 | 4309 | 4323 | 4289 | 4281 | 4300 | 4300.4 | | 3.19 | 3.59 | 3.12 | 3.25 | 3.27 | 3.284 |
| (BT) | 8 | 16 | 4291 | 4258 | 4611 | 4324 | 4299 | 4356.6 | | 4.07 | 3.06 | 9.64 | 4.16 | 3.24 | 4.834 |
| (BT) | 26 | 32 | 4277 | 4266 | 4257 | 4257 | 4216 | 4254.6 | | 3.3 | 3.44 | 3.48 | 3.33 | 3.97 | 3.504 |
| (BT) | 32 | 64 | 4212 | 4314 | 4226 | 4365 | 4208 | 4265 | | 4.7 | 4.18 | 5 | 3.29 | 4.28 | 4.29 |
| | | | | | | | | | | | | | | | |
| | | | | | | | | | | | | | | | |
| | 2 | 4 | 3162 | 3169 | 3163 | 3098 | 3120 | 3142.4 | | 4.8 | 3.88 | 5 | 3.78 | 4.12 | 4.316 |
| | 4 | 8 | 4309 | 4320 | 4290 | 4282 | 4300 | 4300.2 | | 3.19 | 3.59 | 3.12 | 3.25 | 3.27 | 3.284 |
| | 8 | 16 | 3915 | 3937 | 4054 | 3996 | 4185 | 4017.4 | | 4.07 | 3.06 | 9.64 | 4.16 | 3.24 | 4.834 |
| | 16 | 32 | 3906 | 4087 | 3986 | 4031 | 3933 | 3988.6 | | 3.3 | 3.44 | 3.48 | 3.33 | 3.97 | 3.504 |
| | 32 | 64 | 3758 | 3968 | 4026 | 3736 | 3819 | 3861.4 | | 4.7 | 4.18 | 5 | 3.29 | 4.28 | 4.29 |
| | | | | | | | | | | | | | | | |

Part 9 (graph)



Part 10 (graph)



Observation:-- If we compare the average time of threads here. Normal threads are showing better performance than the bounded threads. Here approximately 300 ms difference came between those. This may be due to the way in which mixed is calculated. Unlike chunk a core in mixed calculate on the different arrays. Thus bounding will now lead to more conflicts.

Suppose a thread is running on given cpu core with row k so due to data locality it took the data of k+1 row. If you bound this thread then it have to run on the same core but it will most probably have to execute other row far away than K or K+1. Thus definitely cache miss will be there. But suppose we have not fixed the thread then this thread can go to other core and some other thread can come to this core that will have high probability than previous thread that it will run on K+1 row. This may be the reason why setting the affinity is not helping.

This method is also showing the phenomenon of the saturation when the number of thread is increasing. We can give similar argument that we have given in observation of chunk part.

By doing this experiment I have also observed some unique things that may lead to the data invalidation---

1.when I reboot my PC same algorithm started showing different timing than what It was taking before rebooting.

2.Also when laptop was in battery saver mode it was reducing its processor speed and thus changing the observed time significantly.

Conclusion: --In this experiment I have tried to take every possible aspect that caused this time but at many points I was unable to make correlation between time and cache miss. Also, It is good to measure the interval duration at same time if possible other wise there very high chances that the processor condition change and the process will not give the actual time what it should as scheduler will schedule the things accordingly on number of process available.